

DAY – 3

UNDERSTANDING LARGE LANGUAGE MODELS (LLMs)

On Day 3, we delved deeper into the core of generative AI by understanding Large Language Models (LLMs) — the driving force behind tools like ChatGPT, Claude, Gemini, and others. A Large Language Model is an artificial intelligence model trained on massive datasets to understand and generate human-like language. These models use statistical patterns to perform tasks like translation, summarization, reasoning, and content creation.

We learned that LLMs work by predicting the next word in a sentence, using context and prior knowledge acquired during training.

KEY TERMINOLOGIES

Several important terms were introduced, which are essential to understand the working of LLMs:

- **Token:** A piece of text (word, sub-word, or character) used as input for processing by the model.
- **Parameter:** The weights or values learned during training, which influence how the model responds to a prompt. Modern LLMs have billions of parameters.
- **Prompt:** The input provided to the model to guide its output.
- **Fine-Tuning:** The process of training a pre-trained model further on a specific dataset to tailor its behaviour.
- **Inference:** The phase where a trained model generates output based on new input.

ARCHITECTURE OF A LLM – WHAT'S INSIDE A TRANSFORMER

We studied the inner workings of a Transformer, the foundational architecture behind most LLMs. Key components include:

- **Self-Attention Mechanism:** Allows the model to focus on different words in a sentence relative to each other, enhancing understanding of context.
- **Feed-Forward Neural Networks:** These layers help in transforming the attention-based representations into meaningful output.
- **Positional Encoding:** Since transformers do not process tokens in order, positional encoding provides information about the position of each token in the sequence.

This architecture is the backbone of advanced LLMs like GPT-4, Gemini, and Claude.

TRAINING OF LLMs

We learned about the three key stages involved in training LLMs:

a. Pre-Training:

The initial phase where the model is trained on a vast corpus of general text data. The goal is to learn language structure and general knowledge.

b. Fine-Tuning:

A more focused training phase on specific data or tasks to customize the model's behavior.

c. RLHF (Reinforcement Learning with Human Feedback):

This method involves human reviewers scoring the model's outputs to help it learn preferences, reduce bias, and align better with human values.

APPLICATIONS OF LLMs

LLMs are powering a wide range of applications, such as:

- Virtual assistants and chatbots
- Document summarization
- Code generation and completion
- Language translation
- Creative writing and ideation
- Medical and legal document analysis

LIMITATIONS OF LLMs

While powerful, LLMs come with certain limitations:

- They can generate inaccurate or fabricated information (hallucinations).
- They may not retain context over long conversations.
- High computational requirements make them expensive to train and deploy.
- Their outputs can be sensitive to prompt phrasing.

ETHICAL CONCERNS

We also touched upon the ethical implications of using LLMs:

- **Bias and Fairness:** Models may inherit biases from their training data.
- **Misinformation:** Risk of spreading false information or fake news.
- **Privacy:** Potential to leak personal or sensitive information if trained on unfiltered data.
- **Over-Reliance:** Users may become overly dependent on AI-generated content without verifying accuracy.

CONCLUSION

Day 3 provided an insightful understanding of the technical foundations of LLMs, from their architecture and training to their capabilities and limitations. This session laid a solid theoretical base for working more effectively with AI tools and prompts, helping us become more conscious of both the power and responsibility that come with using such technologies.

